

TIPS FOR WRITING TUTORIAL ESSAYS

Writing a Philosophy Essay

Writing philosophy essays can be difficult, and unlike essays you have written before. It's a skill that takes a lot of practice and correction over time, and you're lucky that you have lots of practice in tutorials to learn how to do it. But we can start with some tips before our first tutorial.

These tips will *not* be exhaustive. They won't cover everything you need to know or every good way to write an essay; but they will provide a starting point. Also a lot of these tips will be only *rules of thumb*; you'll need to get into the practice of writing essays and talking about your essays with your tutor and with each other to think about when these rules are applicable and when they do not. So don't worry if any of them don't work in particular cases!

If you want to read more about essay writing, there are longer guides available online. [Here is one](#) written by the philosopher Jim Pryor; [here is another](#) written by a collection of philosophy tutors at Cambridge.

Three Aims

A good philosophy essay has to do three things:

1. Make the issue comprehensible and compelling.
2. Make an argument.
3. Substantiate that argument.

A good essay will do all of these (actually, a good essay do them all *well*). And we can also think of these three aims in terms of three different *sections* of an essay, sections in each of which you do one of these things. So we'll go through these different sections in turn, and by the end we'll outline a couple of examples of how to structure an essay.

*NB. Notice the importance of **arguments** and **structure***

You should note the focus on *arguments*; essays in other subjects will, of course, include arguments, but their importance is most paramount in a philosophy essay.

I'll also be talking a lot about importance of *structuring* an essay, i.e. how to divide it into sections, what to do at different stages, etc. You should be thinking a lot about how to structure an essay, and you should be self-conscious about this when you're writing, telling the reader what you're going to do, when you've done it, and when you're moving on to doing something else. You also shouldn't worry about writing in continuous prose – giving separate section-titles, and giving indented definitions of views, is very helpful to the reader.

Thinking about structure is also important from a *philosophical* point of view. It may lead you to change how you think about an argument your giving. For example, it may be unclear whether you're arguing against a *premise* in an argument or its *conclusion* – if you force yourself to identify which in the title of a section this can become clearer.

The aim of this focus on structure is to create an essay which is *dialectically sophisticated*. 'Dialectical sophistication' is a concept that's hard to define but easy to understand. It means an argumentative exchange that is sophisticated and goes somewhere. The word 'dialectic' has the same roots as 'dialogue', and the concept is easy to understand with reference to that. Everyone has had annoying discussions with other people, e.g. discussions that don't go anywhere, discussions where your contribution is ignored, discussions that are a series of unrelated points.

Those are dialectically *unsophisticated* discussions. And everyone has had good, productive discussions with other people, e.g. discussions that go somewhere, discussions where the parts relate to one another, discussions in which your points are properly understood. Those are dialectically sophisticated discussions. And even though a philosophy essay is just one voice (yours), it should have the same structure as a good productive conversation, one which features relevant counterarguments, and responses that take into account those counterarguments.

Let's now move on to thinking about the three main aims, and the sections of the essay they will become.

1. Make the issue comprehensible and compelling

This should be the first main section. What it does is give all the background to the issue that you need to in order to give your argument (which we'll get into shortly). In order to this what you have to do is a) describe the issue, b) define terms, and c) briefly describe why there is a compelling issue here. You should be wary of this section being too long though; it shouldn't normally exceed a third of the essay

It is easier to illustrate this with an example, so let's think about how to answer the following question, which I'll refer back to throughout:

Q. "Consequentialism prescribes unconscionable trade-offs." Discuss.

NB. Background for this question

This is (deliberately) not a topic covered in General Philosophy; it's a topic in ethics. Briefly, it's about a problem that *consequentialist* ethical theories face, i.e. theories which claim that what makes an action morally right or morally wrong is the *consequences* of that action. These views are often rejected because they seem to say that it's morally right, in some cases, to make one person's situation worse if it has the 'trade-off' of leading the better consequences overall. A classic example is one which features in the so-called '[trolley problem](#)', originally thought up by [Philippa Foot](#), in which you face the choice of diverting a runaway trolley which will hit five people unless you divert it, but the only alternative direction is one in which the trolley will hit a single person. Most forms of consequentialism say you should divert the train. Many have objected to consequentialism because, they allege, it prescribes making these kinds of trade-offs.




Image: [Jessie Prinz](#)

So in relation to this question, how could you achieve the three things you need to do to make the issue compelling and comprehensible?

a) Describe the issue.

In order to do this, you describe the issue of the question in terms other than those which figure in wording of the question. This makes it clear to reader what the essay is about (and to someone examining ensures you understand what the question is asking). For example, the following captures the issue of the question:

"A problem facing consequentialism is the fact that it prescribes making trade-offs, i.e. that it prescribes

making one person's situation worse if doing so results in better consequences than not."

This captures what is meant by the above question. The following does not:

"A problem facing consequentialism is the fact that it prescribes performing actions that benefit people I've never met, in a 'trade-off' that excludes those who I have a personal relationship with."

The latter is not primarily about trade-offs; it is about a distinct (albeit related) problem consequentialism faces accounting for preference for our nearest and dearest. This may be an *instance* of the trade-off issue, but it doesn't get to the essence of the trade-off issue and gives the reader the impression that the author hasn't fully understood the question.

Therefore, you should try and open with a 1-2 sentence encapsulation of the issue. Don't worry if it's difficult and you don't know how to do it just by reading the question – it is something which emerges out of doing the reading and planning the essay. And there may be more than one way of doing this.

b) Define terms

Student essays are often criticised for not defining terms. But how many terms need defining? Roughly, you need to define anything that won't be understood by someone who hasn't done the reading.

What does that mean in this case. You definitely need to define 'consequentialism' and 'trade-offs' (see that a definition of trade-offs was included in describing the issue). You don't really need to define 'unconscionable' or 'prescribes' – at least not in this essay. 'Unconscionable' here, for example, is not a technical term; it's just a hyperbolic way of saying 'morally wrong'.

When you're reading, look around the literature you're reading for different definitions in order to find the best for the essay; the first definition you find might be wrong, too complex for your purposes, etc.

NB. A couple of ways of going wrong with defining terms

1. Including needless extra information; being too specific.

Bad: "Consequentialism is the view that the moral rightness of an action depends on the pleasure and pain brought about by the action."

Defining consequentialism in terms of pleasure and pain is too specific. Other forms of consequentialism don't hold that pleasure and pain are the fundamental consequences we should bring about; others claim we should make people happy in a richer sense, rather than merely increasing their pleasure and decreasing their pain. More importantly, it's overly specific in a way that's irrelevant to the issue of trade-offs. Versions of consequentialism which claim that we should make people happy in a richer sense also face issue with trade-offs.

Better: "Consequentialism is the view that the moral rightness of an action depends on the goodness or badness of the consequences brought about by that action."

However, this is still not ideal for this purposes of *this* essay, because it is in another respect too *unspecific*. In order to work out whether consequentialism prescribes trade-offs we will need to specify *how* rightness depends on the goodness or badness of consequences. A natural starting point is a view that many consequentialists defend: *Maximizing* consequentialism.

Best: "Maximizing consequentialism is the view that an action is morally right if and only if, out of all the options available to the agent, it brings about the consequences with the highest balance of goodness over badness."

However, you should note that this is a *specific version* of consequentialism, and that answering the question fully may involve examining other forms of consequentialism which understand the way in which rightness depends on consequences differently (more on this later).

2. Defining terms in way that makes the question uninteresting; too easy to answer.

Bad: “A trade-off is doing something morally wrong to someone, when that has consequences that are better for everyone overall.”

This is bad because if a trade-off is defined as morally wrong. This means that if we accept that consequentialism prescribes trade-offs, then consequentialism is false by definition. This is an instance of what’s called ‘arguing against a straw man’, i.e. arguing against an artificially implausible position.

Better: “A trade-off is making one person’s situation worse, when that has consequences that are better for everyone overall.”

This term is also one which it is good to illustrate with an *example*. You could use Foot’s trolley case, but there are many others in the literature; you could also come up with your own.

NB. When taking examples from the literature, you should be wary of including extraneous material. For example, Foot’s trolley case was *not* originally used as an instance of a trade-off *as such*, but appealed to when investigating the ethical difference between *doing* and *allowing* harm.

c) Make the issue compelling

What you now need to do is to say why there is an issue here. This means making the question compelling enough to warrant an essay talking about it. In this essay, you need to explain why trade-offs create a problem that the defenders of consequentialism at least need to answer (you don’t need to claim that it *can’t* be answered – you may in the end defend consequentialism).

How do you do this? For some essays, if you do the previous two things well you may have already done it. For others, it may involve a bit more than describing the issue and defining terms, why the issue is one that requires a response. View it from as if you’re trying to persuade someone, who hasn’t read the material, to spend more time on considering whether this problem can be answered. In this case, it may just involve highlighting that making trade-offs doesn’t look like what we would normally regard as morally right action, and so if consequentialism does prescribe them, that would conflict with our common sense understanding of morality.

2. Make your argument

This is the centrepiece of the essay. Here you have to make an argument, specifically an argument which *answers the essay question*. Be very careful to make sure that your essays answer the question you’re given. For example, only discussing *other* objections to consequentialism won’t answer this question; defending consequentialism in general, but not answering the problem of trade-offs won’t answer this question.

That needn’t be constricting; there are many different ways to answer a question like the one above. For simplicity’s sake, however, we’ll concentrate on two broad ways of answering it as illustrative examples:

- a) *A Positive Answer*, i.e. an essay arguing that *yes*, consequentialism *does* prescribe unconscionable trade-offs, and *should* be rejected for that reason.
- b) *A Negative Answer*, i.e. an essay arguing that *no*, consequentialism does *not* prescribe unconscionable trade-offs, and should *not* be rejected for that reason.

These kinds of answer are available because of the kind of question we’re answering. Because it’s a question which is asking you to assess an argument against a view, you can either conclude that that argument is *good* or you can conclude that it’s *bad*. Therefore, focusing on these two kinds of answer is helpful because it will be applicable to any question that asks you to assess an argument against a

position. However, they have broader application, because they really apply whenever there's a choice between *defending* or *rejecting* a claim or view or argument.

NB. A reminder that this is not exhaustive.

Here, more than ever, these recommendations are not exhaustive. Many arguments won't fall into either these categories. A lot will though, and thinking about how to formulate arguments of these two categories will help when you come to formulate different kinds of arguments.

Let's now move on to how to do this in relation to our example, with each of these two options.

a) Positive Answer

What if you're arguing that consequentialism should be rejected because it prescribes trade-offs? If you do this, you have to defend an argument like the one suggested by our question. We can divide this task into two parts: i) stating an argument, ii) defending that argument.

i) Stating the argument

Doing this involves about what your argument is. One simple but effective way to do this is to formulate the argument in premises and a conclusion.

For example, with this essay we might have the following:

Argument from Trade-Offs

1. Consequentialism prescribes making one person's situation worse if that has better consequences overall than not doing so.
2. It's morally wrong to make one person's situation worse even if it has better consequences overall than not doing so.
- C. Consequentialism prescribes doing what's morally wrong.

This captures the argument suggested by the question. Premise 1 formulates the idea that consequentialism prescribes trade-offs, and premise 2 formulates the idea that trade-offs are unconscionable. And since consequentialism is, as our definition above said, a theory about morally right and wrong action, then if it prescribe doing what's morally wrong, it's a false theory.

NB. But do you always need to formulate arguments in premises and conclusions?

The quick answer is no. Sometimes the argument will be too obvious, e.g. if it's just a counterexample to a theory, then formulating an argument may be tedious; other times you'll be able to state the argument clearly enough in continuous prose.

However, at this stage, I'd suggest you try and put arguments premises and conclusions whenever possible. It's a very good exercise, because it forces you to get clear about what exactly you are arguing for (or arguing against), and it really helps when thinking about objections to the argument (more of this very shortly). Though don't worry if your initial attempts at formalising arguments results in ones we end up criticising and correcting in the tutorial; that's just part of the process of learning how to formulate arguments.

ii) Defending that argument

Having set out the argument clearly, the second main thing you have to do in this section is to *defend* this argument. What might that involve in this case? This is a lot easier now you've got the argument

formulated in premises and conclusions. What you have to do is go through each premise and explain why they are intuitively compelling. Let's do this here.

1. Consequentialism prescribes making one person's situation worse if that has better consequences overall than not doing so.

How one might argue that this premise is plausible? A good way in this case would be to do this would be a) give your definition of consequentialism, and b) give (or refer back to) an example in which it plausibly prescribe making one person's situation worse because that has better consequences overall. Let's try and do this with our definition of consequentialism:

Maximizing consequentialism: An action is morally right if and only if, out of all the options available, it brings about the consequences with the highest balance of goodness over badness.

We now just have to argue that this version of consequentialism prescribes making one person's situation worse. In this case, the best way to do this would be to outline an example or thought experiment in which the action that brings about the consequences with the highest balance of goodness over badness does in fact involve making someone's situation worse. In our case, outlining how it does so in Philippa Foot's trolley case would likely be sufficient.

Let's move on to the next premise.

2. It's morally wrong to make one person's situation worse even if it has better consequences overall than not doing so.

How one might argue that this premise is plausible?

What you shouldn't really do here is defend premise 2 on the basis that it follows from *some other moral theory* (such as Kantianism). This is because you need to give *theory-neutral* reasons to reject consequentialism.

One way of arguing of arguing for premise 2 would be to argue that it is an expression of various important moral concepts we appeal to and employ in our moral evaluations of people's conduct, concepts such as rights, harm, autonomy, freedom, and innocence. You could argue that these concepts are an important part of our moral evaluation of people and that for them to be so premise 2 must be true. *Alternatively*, you could argue that the trade-off examples and thought experiments illicit in us strong intuition that it's morally wrong to make one person's situation worse for the benefit of others, and that this intuition provides reason to think that premise 2 is true.

These are only examples of ways of arguing for premise 2, and they would have to be more developed. But they are examples of attempts to provide a *theory neutral* justification of premise 2 – i.e. a justification which would hold independently of any particular moral theory.

And now made each of the two premises of our argument compelling, we can now claim we have a compelling argument with the following conclusion:

- C. Consequentialism prescribes doing what's morally wrong.

NB. The question of the argument's validity

Before you draw the conclusion, you actually need to do more than defend the premises, you also need to ensure that the argument is 'valid', i.e. that the conclusion actually does follow from the

premises. A *valid* argument is one where if the premises are true, then the conclusion can't *possibly* fail to be true. You'll learn about validity in logic, so once you've had a few logic classes, you'll be able to understand what a valid argument is and how to make one.

Making an argument with intuitively compelling premises, as above, is really all you need to do in this central section of the essay. All you need to do is not defend the premises against every counterargument here; that's for the section in which you *substantiate* your argument.

b) Negative Answer

A lot of the things we'd said about giving a *positive* answer to the question are also true of giving a *negative* answer. The key *difference* in writing a negative answer, you're arguing that some argument *fails*. Given this, it's extremely helpful to formulate the argument you are arguing *against*, so you can specify the exact reason why the argument fails. If you have an argument with premises and a conclusion, you can:

- 1) Say *which* premise you'll be rejecting.
- 2) Give *reasons* for rejecting that premise.

With our consequentialism question, we can just use the argument we formulated for the positive answer for this purpose:

Argument from Trade-Offs

1. Consequentialism prescribes making one person's situation worse if that has better consequences overall than not doing so.
2. It's morally wrong to make one person's situation worse even if it has better consequences overall than not doing so.
- C. Consequentialism prescribes doing what's morally wrong.

Given this argument has two premises, this gives (at least) two possible directions a negative answer could go. In other words, you could:

Reject premise 1: You could argue that we don't have good reason to think that consequentialism prescribes trade-offs. *One* way of doing this would be to argue that it's actually false of Maximizing consequentialism that it prescribes trade-offs. Perhaps you might say it's only in artificial situations that never actually happen that Maximizing consequentialism prescribes trade-offs. *Alternatively*, you could argue that some alternative version of consequentialism, such as *Rule* consequentialism, doesn't in fact prescribe making trade-offs

Reject premise 2: You could argue that trade-offs are not in fact unconscionable. You might say that the only reason we find them unconscionable is because a feeling of squeamishness at the prospect of causing someone's death in the thought experiments, and that a feeling of squeamishness is not relevant to whether or not the acts in question are morally wrong.

How much time do you have to spend arguing against the premise you choose to reject? Again, as with the positive answer, all you need to do in this section is make your rejection of the premises intuitively compelling. You don't need to argue against all possible counterarguments to your rejection of this premise; that's for the next section.

NB. *What about arguing against both premises?*

You may be tempted to argue against both. This is fine; but space is at a premium in an essay like this. Given that an argument fails if just *one* of its premises is false, it's better to give a stronger, more sustained case against one premise (I'll explain why in more detail in the next section).

3. Substantiate your argument

What do I mean by saying you need to substantiate your argument? An unsubstantiated argument is, roughly, an argument which you make but you don't defend it against objections, especially obvious objections. Perhaps, you give an argument with a premise which everyone always has the same objection to, but you don't answer that response. Even if you think that objection is terrible, and there's a very easy answer to it, unless you give that answer, your argument has not been substantiated. To substantiate this argument, you need to answer that objection.

So substantiating an argument involves responding to possible counterarguments. But, as I've said, you'll have limited space in the essays you'll be writing. So how many counterarguments should you discuss? And if you can only discuss some counterarguments, which ones should you focus on? There are more dialectically sophisticated ways to do this, and a few aims can help here.

1. Focus on *stronger* responses

You should try and focus on the strongest counterarguments to the argument you make. There's nothing worse than someone rejecting an extremely implausible objection, and then declaring that they've defended their view against criticism (this is another example of arguing against a 'straw man'). Therefore, you should focus on the strongest, most immediate, and most obvious responses to the arguments you make.

How do you tell which counterarguments are the strongest? There's no simple answer to this. In the end you are going to have to make a judgement call about which counterarguments you see as most needing an answer. One tip I suggest though is that you base it on *both* your own judgement, and on which response which features most prominently in the reading, your lectures, and discussion with your peers.

2. Focus on ways the view being argued against could be *modified* in response to your argument

It's unsatisfying if you object to a view, say, by giving a counterexample, but don't then consider ways in which that view might be modified so as to deal with your counterexample. This is especially unsatisfying if there is an *obvious* way to modify the view. This is yet another instance of arguing against a 'straw man'. For example, a key move of this kind in relation to our question about trade-offs is the move from *Maximizing* consequentialism to *Rule* consequentialism:

Rule consequentialism: An action is morally right if and only if, out of all the options available, it would be part of a set of rules which if universally followed would bring about the consequences with the highest balance of goodness over badness.

Rule consequentialism does not, its proponents argue, prescribe trade-offs, because although a trade-off might have better consequences in some cases, a *rule* prescribing making trade-offs would not have the best consequences.

However, discussing modified views which deal with your initial counterexample does not mean your initial objection wasn't a good one. Your initial counterexample can still play an important dialectical role *if* you have good arguments against the modified views. We can outline a couple of ways of doing this:

If the modified views still face similarly structured counterexamples, then you can show that the modification didn't really deal with the essence of your counterexample. With our question and the Rule consequentialist alternative, you could do this by coming up with a case in which Rule consequentialism would prescribe making a trade-off. This might be the case in an imagined world in which the options of making trade-offs with better overall consequences were much

more prevalent than in our world, so a rule prescribing making them would plausibly have better consequences.

If the modified views are independently implausible or unmotivated, then you can show that modifying the view in response to your counterexample leads to an implausible view. With the Rule consequentialist alternative, you could do this by arguing that the requirement that morally right actions form part of a set of rules is implausible. One might argue that it is implausible as a version of consequentialism because it seems that breaking the rules in some cases would produce the best consequences.

This is just one example, but it's a process you should try and carry out where possible and relevant; it's especially useful in essays in which you try and *analyse* a concept and consider possible *counterexamples* to the analyses. But more generally, discussing possible modifications of the view you're objecting to, and *then* arguing that even the modified views are still vulnerable to the same objection, results in a more substantiated argument and a much more dialectically sophisticated essay

3. Focus on only as many responses as you *need to* in order to make your point

This aim is harder to state, but really helps with structuring an essay in a dialectically sophisticated way. You should only focus on those counterarguments which you need to answer in order to make your point, and you should ignore discussing responses or giving arguments that are not relevant to that. We can illustrate this in relation to our positive and negative kinds of answer.

If you're giving a positive answer – i.e. if you're *defending* an argument or a view – then if there is *any* good counterargument to your view, your defence of that argument is undermined. Therefore, if you're giving a positive answer in response to our question about consequentialism, if either of the two premises is false, the whole argument is undermined. Therefore, it's not enough to defend that argument *just* to defend premise 1, i.e. *just* to defend the claim that consequentialism prescribes trade-offs. You also need to defend premise 2, i.e. to defend the claim that making trade-offs is morally wrong. Therefore, you need to respond to possible counterarguments to *both* premises. If there are counterarguments available against one of the premises, it's not enough to just defend the other premise against possible counterarguments, however well you do that.

On the other hand, if you're giving a negative answer – i.e. if you're *rejecting* an argument – if you give one *good* response to that argument, that *by itself* provides sufficient reason to reject it. If you already make one good objection, giving *another* reason to reject the argument is surplus to requirements. However, once you give your response, you're now *making your own argument*, and if any step of *your* response fails, then your whole response is undermined. Therefore, if you're giving a negative answer to the question, it's more important to discuss and respond to possible objections to your counterargument, rather than considering further possible counterarguments to the main argument.

With these two points in mind, that gives us two possible structures depending on what kind of answer you'll be giving to the essay question:

Example of a Positive Answer

1. Introduction
2. Background
3. Main argument (which you'll be defending)
4. 1st Objection to main argument, and your response
5. 2nd Objection to main argument, and your response
6. Conclusion

Example of Negative Answer

1. Introduction
2. Background
3. Main argument (which you'll be rejecting)
4. Your counterargument to main argument
5. 1st response to your counterargument, and your reply
6. 2nd response to your counterargument, and your reply
7. Conclusion

NB. What about introductions and conclusions?

I've included introductions and conclusions in both of these structures, but I haven't said what they should include. An introduction should only really be a prospective summary of what you're going to do (so best written last), though it should also include a brief description of the issue of the kind we mentioned earlier. A conclusion really only needs to be a retrospective summary of what you have done.

You *don't* need to include, in either of these two sections, any grand statements about the importance of the issue and its place in history. These kinds of statements are hard to substantiate, and they give essays a clichéd feel; just get straight into the business of giving the arguments.

These are two ways to structure a dialectically sophisticated philosophy essay. However, they are only *examples* of how to structure an essay, examples which might work for one question, but won't work for others. The point of these two structures is so that you think about the structure of your essay and think about, e.g., whether it's missing any steps, or whether any steps are unnecessary.

More generally, the point of all my suggestions in this guide is to encourage you to be self-conscious about the structure of your essays and the roles of different arguments within them. If you think self-consciously about what you exactly *doing* at each stage of your essay, you'll coming up with dialectically sophisticated philosophy essays of your own in no time. Don't worry if you find this hard at first. It is hard, especially when you're not working with constructed examples like I have been. But if you persevere, you even might end up enjoying this aspect of writing!

Good Luck!

NB. Do get in touch

This guide is a work-in-progress, so please do get in touch if anything is unclear, or if you have any feedback about what it would be good to spend more or less time on. I'd very much appreciate it!

Alexander Greenberg
alexander.greenberg@lmh.ox.ac.uk